

## sklearn.metrics.precision\_recall\_fscore\_support

»

```
sklearn.metrics.precision_recall_fscore_support(y_true, y_pred, beta=1.0, labels=None,  
pos_label=1, average=None, warn_for=('precision', 'recall', 'f-score'), sample_weight=None)
```

[\[source\]](#)

Compute precision, recall, F-measure and support for each class

The precision is the ratio  $tp / (tp + fp)$  where  $tp$  is the number of true positives and  $fp$  the number of false positives. The precision is intuitively the ability of the classifier not to label as positive a sample that is negative.

The recall is the ratio  $tp / (tp + fn)$  where  $tp$  is the number of true positives and  $fn$  the number of false negatives. The recall is intuitively the ability of the classifier to find all the positive samples.

The F-beta score can be interpreted as a weighted harmonic mean of the precision and recall, where an F-beta score reaches its best value at 1 and worst score at 0.

The F-beta score weights recall more than precision by a factor of  $\beta$ .  $\beta == 1.0$  means recall and precision are equally important.

The support is the number of occurrences of each class in  $y\_true$ .

If `pos_label` is `None` and in binary classification, this function returns the average precision, recall and F-measure if `average` is one of 'micro', 'macro', 'weighted' or 'samples'.

Read more in the [User Guide](#).

---

**Parameters:** **y\_true** : 1d array-like, or label indicator array / sparse matrix

Ground truth (correct) target values.

**y\_pred** : 1d array-like, or label indicator array / sparse matrix

Estimated targets as returned by a classifier.

**beta** : float, 1.0 by default

The strength of recall versus precision in the F-score.

**labels** : list, optional

The set of labels to include when `average != 'binary'`, and their order if `average is None`. Labels present in the data can be excluded, for example

to calculate a multiclass average ignoring a majority negative class, while labels not present in the data will result in 0 components in a macro average. For multilabel targets, labels are column indices. By default, all labels in `y_true` and `y_pred` are used in sorted order.

**pos\_label** : str or int, 1 by default

The class to report if `average='binary'`. Until version 0.18 it is necessary to set `pos_label=None` if seeking to use another averaging method over binary targets.

**average** : string, [None (default), 'binary', 'micro', 'macro', 'samples', 'weighted']

If None, the scores for each class are returned. Otherwise, this determines the type of averaging performed on the data:

'binary':

Only report results for the class specified by `pos_label`. This is applicable only if targets (`y_{true,pred}`) are binary.

'micro':

Calculate metrics globally by counting the total true positives, false negatives and false positives.

'macro':

Calculate metrics for each label, and find their unweighted mean. This does not take label imbalance into account.

'weighted':

Calculate metrics for each label, and find their average, weighted by support (the number of true instances for each label). This alters 'macro' to account for label imbalance; it can result in an F-score that is not between precision and recall.

'samples':

Calculate metrics for each instance, and find their average (only meaningful for multilabel classification where this differs from [accuracy\\_score](#)).

Note that if `pos_label` is given in binary classification with `average != 'binary'`, only that positive class is reported. This behavior is deprecated and will change in version 0.18.

**warn\_for** : tuple or set, for internal use

This determines which warnings will be made in the case that this function is being used to return only one of its metrics.

**sample\_weight** : array-like of shape = [n\_samples], optional

Sample weights.

**Returns:**      **precision:** float (if average is not None) or array of float, shape = [n\_unique\_labels] :

**recall:** float (if average is not None) or array of float, , shape = [n\_unique\_labels] :

**fbeta\_score:** float (if average is not None) or array of float, shape = [n\_unique\_labels] :

**support:** int (if average is not None) or array of int, shape = [n\_unique\_labels] :

The number of occurrences of each label in y\_true.

## References

[R175] [Wikipedia entry for the Precision and recall](#)

[R176] [Wikipedia entry for the F1-score](#)

[R177] *Discriminative Methods for Multi-labeled Classification Advances in Knowledge Discovery and Data Mining (2004), pp. 22-30 by Shantanu Godbole, Sunita Sarawagi*  
<<http://www.godbole.net/shantanu/pubs/multilabelsvm-pakdd04.pdf>>

## Examples

```
>>> from sklearn.metrics import precision_recall_fscore_support
>>> y_true = np.array(['cat', 'dog', 'pig', 'cat', 'dog', 'pig'])
>>> y_pred = np.array(['cat', 'pig', 'dog', 'cat', 'cat', 'dog'])
>>> precision_recall_fscore_support(y_true, y_pred, average='macro')
...
(0.22..., 0.33..., 0.26..., None)
>>> precision_recall_fscore_support(y_true, y_pred, average='micro')
...
(0.33..., 0.33..., 0.33..., None)
>>> precision_recall_fscore_support(y_true, y_pred, average='weighted')
...
(0.22..., 0.33..., 0.26..., None)
```

It is possible to compute per-label precisions, recalls, F1-scores and supports instead of averaging: >>> precision\_recall\_fscore\_support(y\_true, y\_pred, average=None, ... labels=['pig', 'dog', 'cat']) ... # doctest: +ELLIPSIS, +NORMALIZE\_WHITESPACE (array([ 0. , 0. , 0.66...]),

array([ 0., 0., 1.]), array([ 0. , 0. , 0.8]), array([2, 2, 2]))