

## sklearn.neighbors.KNeighborsClassifier

» `class sklearn.neighbors.KNeighborsClassifier(n_neighbors=5, weights='uniform', algorithm='auto', leaf_size=30, p=2, metric='minkowski', metric_params=None, n_jobs=1, **kwargs)` [\[source\]](#)

Classifier implementing the k-nearest neighbors vote.

Read more in the [User Guide](#).

**Parameters:** **n\_neighbors** : int, optional (default = 5)

Number of neighbors to use by default for `k_neighbors` queries.

**weights** : str or callable

weight function used in prediction. Possible values:

- 'uniform' : uniform weights. All points in each neighborhood are weighted equally.
- 'distance' : weight points by the inverse of their distance. in this case, closer neighbors of a query point will have a greater influence than neighbors which are further away.
- [callable] : a user-defined function which accepts an array of distances, and returns an array of the same shape containing the weights.

Uniform weights are used by default.

**algorithm** : {'auto', 'ball\_tree', 'kd\_tree', 'brute'}, optional

Algorithm used to compute the nearest neighbors:

- 'ball\_tree' will use [BallTree](#)
- 'kd\_tree' will use [KDTree](#)
- 'brute' will use a brute-force search.
- 'auto' will attempt to decide the most appropriate algorithm based on the values passed to [fit](#) method.

Note: fitting on sparse input will override the setting of this parameter, using brute force.

**leaf\_size** : int, optional (default = 30)

Leaf size passed to [BallTree](#) or [KDTree](#). This can affect the speed of the construction and query, as well as the memory required to store the tree.

The optimal value depends on the nature of the problem.

**metric** : string or DistanceMetric object (default = 'minkowski')

the distance metric to use for the tree. The default metric is minkowski, and with  $p=2$  is equivalent to the standard Euclidean metric. See the documentation of the DistanceMetric class for a list of available metrics.

**p** : integer, optional (default = 2)

Power parameter for the Minkowski metric. When  $p = 1$ , this is equivalent to using `manhattan_distance` (l1), and `euclidean_distance` (l2) for  $p = 2$ . For arbitrary  $p$ , `minkowski_distance` (l\_p) is used.

**metric\_params** : dict, optional (default = None)

Additional keyword arguments for the metric function.

**n\_jobs** : int, optional (default = 1)

The number of parallel jobs to run for neighbors search. If -1, then the number of jobs is set to the number of CPU cores. Doesn't affect `fit` method.

**See also:** [RadiusNeighborsClassifier](#), [KNeighborsRegressor](#), [RadiusNeighborsRegressor](#), [NearestNeighbors](#)

## Notes

See [Nearest Neighbors](#) in the online documentation for a discussion of the choice of `algorithm` and `leaf_size`.

**Warning:** Regarding the Nearest Neighbors algorithms, if it is found that two neighbors, neighbor  $k+1$  and  $k$ , have identical distances but different labels, the results will depend on the ordering of the training data.

[http://en.wikipedia.org/wiki/K-nearest\\_neighbor\\_algorithm](http://en.wikipedia.org/wiki/K-nearest_neighbor_algorithm)

## Examples

```
>>> X = [[0], [1], [2], [3]]
>>> y = [0, 0, 1, 1]
>>> from sklearn.neighbors import KNeighborsClassifier
>>> neigh = KNeighborsClassifier(n_neighbors=3)
>>> neigh.fit(X, y)
KNeighborsClassifier(...)
>>> print(neigh.predict([[1.1]]))
[0]
>>> print(neigh.predict_proba([[0.9]]))
[[ 0.66666667  0.33333333]]
```

>>>

## Methods

<code>fit(X, y)</code>	Fit the model using X as training data and y as target values
<code>get_params([deep])</code>	Get parameters for this estimator.
<code>kneighbors([X, n_neighbors, return_distance])</code>	Finds the K-neighbors of a point.
<code>kneighbors_graph([X, n_neighbors, mode])</code>	Computes the (weighted) graph of k-Neighbors for points in X
<code>predict(X)</code>	Predict the class labels for the provided data
<code>predict_proba(X)</code>	Return probability estimates for the test data X.
<code>score(X, y[, sample_weight])</code>	Returns the mean accuracy on the given test data and labels.
<code>set_params(**params)</code>	Set the parameters of this estimator.

`__init__(n_neighbors=5, weights='uniform', algorithm='auto', leaf_size=30, p=2, metric='minkowski', metric_params=None, n_jobs=1, **kwargs)` [\[source\]](#)

`fit(X, y)` [\[source\]](#)

Fit the model using X as training data and y as target values

**Parameters:** **X** : {array-like, sparse matrix, BallTree, KDTree}

Training data. If array or matrix, shape [n\_samples, n\_features], or [n\_samples, n\_samples] if metric='precomputed'.

**y** : {array-like, sparse matrix}

Target values of shape = [n\_samples] or [n\_samples, n\_outputs]

`get_params(deep=True)` [\[source\]](#)

Get parameters for this estimator.

**Parameters:** **deep**: boolean, optional :

If True, will return the parameters for this estimator and contained subobjects that are estimators.

**Returns:** **params** : mapping of string to any

Parameter names mapped to their values.

`kneighbors(X=None, n_neighbors=None, return_distance=True)` [\[source\]](#)

Finds the K-neighbors of a point.

Returns indices of and distances to the neighbors of each point.

**Parameters:** **X** : array-like, shape (n\_query, n\_features), or (n\_query, n\_indexed) if metric == 'precomputed'

The query point or points. If not provided, neighbors of each indexed point are returned. In this case, the query point is not considered its own neighbor.

**n\_neighbors** : int

Number of neighbors to get (default is the value passed to the constructor).

**return\_distance** : boolean, optional. Defaults to True.

If False, distances will not be returned

**Returns:** **dist** : array

Array representing the lengths to points, only present if return\_distance=True

**ind** : array

Indices of the nearest points in the population matrix.

## Examples

In the following example, we construct a NeighborsClassifier class from an array representing our data set and ask who's the closest point to [1,1,1]

```
>>> samples = [[0., 0., 0.], [0., .5, 0.], [1., 1., .5]]
>>> from sklearn.neighbors import NearestNeighbors
>>> neigh = NearestNeighbors(n_neighbors=1)
>>> neigh.fit(samples)
NearestNeighbors(algorithm='auto', leaf_size=30, ...)
>>> print(neigh.kneighbors([[1., 1., 1.]])
(array([[ 0.5]]), array([[2]]...))
```

As you can see, it returns [[0.5]], and [[2]], which means that the element is at distance 0.5 and is the third element of samples (indexes start at 0). You can also query for multiple points:

```
>>> X = [[0., 1., 0.], [1., 0., 1.]]
>>> neigh.kneighbors(X, return_distance=False)
array([[1],
       [2]]...)
```

`kneighbors_graph(X=None, n_neighbors=None, mode='connectivity')`

[\[source\]](#)

---

Computes the (weighted) graph of k-Neighbors for points in X

---

**Parameters:** **X** : array-like, shape (n\_query, n\_features), or (n\_query, n\_indexed) if metric == 'precomputed'

The query point or points. If not provided, neighbors of each indexed point are returned. In this case, the query point is not considered its own neighbor.

**n\_neighbors** : int

Number of neighbors for each sample. (default is value passed to the constructor).

**mode** : {'connectivity', 'distance'}, optional

Type of returned matrix: 'connectivity' will return the connectivity matrix with ones and zeros, in 'distance' the edges are Euclidean distance between points.

---

**Returns:** **A** : sparse matrix in CSR format, shape = [n\_samples, n\_samples\_fit]

n\_samples\_fit is the number of samples in the fitted data A[i, j] is assigned the weight of edge that connects i to j.

---

**See also:** [NearestNeighbors.radius\\_neighbors\\_graph](#)

## Examples

```
>>> X = [[0], [3], [1]]
>>> from sklearn.neighbors import NearestNeighbors
>>> neigh = NearestNeighbors(n_neighbors=2)
>>> neigh.fit(X)
NearestNeighbors(algorithm='auto', leaf_size=30, ...)
>>> A = neigh.kneighbors_graph(X)
>>> A.toarray()
array([[ 1.,  0.,  1.],
       [ 0.,  1.,  1.],
       [ 1.,  0.,  1.]])
```

>>>

**predict(X)**

[\[source\]](#)

Predict the class labels for the provided data

---

**Parameters:** **X** : array-like, shape (n\_query, n\_features), or (n\_query, n\_indexed) if metric == 'precomputed'

Test samples.

---

**Returns:** **y** : array of shape [n\_samples] or [n\_samples, n\_outputs]

Class labels for each data sample.

---

`predict_proba(X)`

[\[source\]](#)

Return probability estimates for the test data X.

**Parameters:** **X** : array-like, shape (n\_query, n\_features), or (n\_query, n\_indexed) if metric == 'precomputed'

Test samples.

---

**Returns:** **p** : array of shape = [n\_samples, n\_classes], or a list of n\_outputs

of such arrays if n\_outputs > 1. The class probabilities of the input samples. Classes are ordered by lexicographic order.

---

`score(X, y, sample_weight=None)`

[\[source\]](#)

Returns the mean accuracy on the given test data and labels.

In multi-label classification, this is the subset accuracy which is a harsh metric since you require for each sample that each label set be correctly predicted.

---

**Parameters:** **X** : array-like, shape = (n\_samples, n\_features)

Test samples.

**y** : array-like, shape = (n\_samples) or (n\_samples, n\_outputs)

True labels for X.

**sample\_weight** : array-like, shape = [n\_samples], optional

Sample weights.

---

**Returns:** **score** : float

Mean accuracy of self.predict(X) wrt. y.

---

`set_params(**params)`

[\[source\]](#)

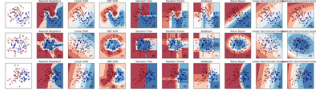
Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as pipelines). The former have parameters of the form `<component>__<parameter>` so that it's possible to update each component of a nested object.

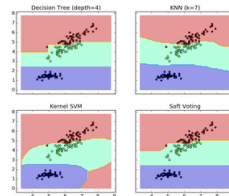
**Returns:** `self` :

## Examples using `sklearn.neighbors.KNeighborsClassifier`

»



Classifier comparison



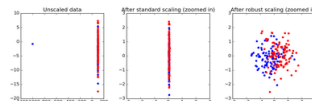
Plot the decision boundaries of a VotingClassifier



Digits Classification Exercise



Nearest Neighbors Classification



Robust Scaling on Toy Data



Classification of text documents using sparse features